

1. Dropout on Linear Regression Recall that linear regression optimizes:

$$\mathcal{L}(\mathbf{w}) = \|\mathbf{y} - X\mathbf{w}\|_2^2 \quad (1)$$

One way of using *dropout* on the d -dimensional input features \mathbf{x}_i involves keeping each feature at random with probability p (and zeroing it out if not kept). This makes our learning objective effectively become

$$\mathcal{L}(\tilde{\mathbf{w}}) = \mathbb{E}_{R \sim \text{Bernoulli}(p)} \left[\|\mathbf{y} - (R \odot X)\tilde{\mathbf{w}}\|_2^2 \right] \quad (2)$$

where \odot is the element-wise product, and the random binary matrix $R \in \{0, 1\}^{n \times d}$ is such that $R_{i,j} \sim_{i.i.d} \text{Bernoulli}(p)$. We use $\tilde{\mathbf{w}}$ to remind you that this is learned by dropout.

Show that we can manipulate (2) to eliminate the expectations and get:

$$\mathcal{L}(\tilde{\mathbf{w}}) = \|\mathbf{y} - pX\tilde{\mathbf{w}}\|_2^2 + p(1-p)\|\tilde{\Gamma}\tilde{\mathbf{w}}\|_2^2 \quad (3)$$

with $\tilde{\Gamma}$ being a diagonal matrix whose j -th diagonal entry is the norm of the j -th column of the training matrix X .

Solution: Let $P = R \odot X$ where \odot is the element-wise multiplication. Therefore, we have:

$$\|\mathbf{y} - P\mathbf{w}\|_2^2 = \mathbf{y}^T \mathbf{y} - 2\mathbf{w}^T P^T \mathbf{y} + \mathbf{w}^T P^T P \mathbf{w} \quad (4)$$

That is:

$$\mathbb{E}_{R \sim \text{Bernoulli}(p)} [\|\mathbf{y} - R \odot X \mathbf{w}\|_2^2] = \mathbb{E}_R [\mathbf{y}^T \mathbf{y} - 2\mathbf{w}^T P^T \mathbf{y} + \mathbf{w}^T P^T P \mathbf{w}] \quad (5)$$

Since the expected value of a matrix is the matrix of the expected value of its elements, we have that

$$\mathbb{E}_R [P]_{ij} = \mathbb{E}_R [(R \odot X)_{ij}] = X_{ij} \mathbb{E}_R [R_{ij}] = pX_{ij} \quad (6)$$

It follows that:

$$\mathbb{E}_R [2\mathbf{w}^T P^T \mathbf{y}] = 2p\mathbf{w}^T X^T \mathbf{y} \quad (7)$$

and:

$$(\mathbb{E}_R [P^T P])_{ij} = \sum_{k=1}^N \mathbb{E}_R [R_{ki} R_{kj} X_{ki} X_{kj}] \quad (8)$$

where:

$$\mathbb{E}_R[(P^T P)]_{ij} = \begin{cases} \sum_{k=1}^N \mathbb{E}_R[R_{ki} R_{kj} X_{ki} X_{kj}] = \sum_{k=1}^N \mathbb{E}_R[R_{ki}] \mathbb{E}_R[R_{kj}] X_{ki} X_{kj} = p^2 (X^T X)_{ij} & \text{if } i \neq j \\ \sum_{k=1}^N \mathbb{E}_R[R_{ki}^2 X_{ki} X_{kj}] = \sum_{k=1}^N \mathbb{E}_R[R_{ki}^2] X_{ki} X_{kj} = p (X^T X)_{ij} & \text{if } i = j \end{cases} \quad (9)$$

Finally, we note that :

$$(\mathbb{E}_R[(P^T P)])_{ij} - p^2 (X^T X)_{ij} = \begin{cases} 0 & \text{if } i \neq j \\ (p - p^2) (X^T X)_{ij} & \text{if } i = j \end{cases} \quad (10)$$

we now can put everything together as follow:

$$\mathcal{L}(w) = \mathbb{E}_R[\|y - R \odot Xw\|_2^2] \quad (11)$$

$$= \mathbb{E}_R[y^T y - 2w^T P^T y + w^T P^T P w] \quad (12)$$

$$= y^T y - 2pw^T X^T y + p^2 w^T X^T X w - p^2 w^T X^T X w + w^T \mathbb{E}_R[P^T P] w \quad (13)$$

$$= \|y - pXw\|_2^2 + (w^T \mathbb{E}_R[P^T P] w - p^2 w^T X^T X w) \quad (14)$$

$$= \|y - pXw\|_2^2 + w^T (\mathbb{E}_R[P^T P] w - p^2 w) \quad (15)$$

$$= \|y - pXw\|_2^2 + (p^2 - p) w^T (\text{diag}(X^T X)) w \quad (16)$$

$$= \|y - pXw\|_2^2 + p(1 - p) w^T (\text{diag}(X^T X)) w \quad (17)$$

$$= \|y - pXw\|_2^2 + p(1 - p) \|\tilde{\Gamma} w\|_2^2 \quad (18)$$

$$(19)$$

where $\text{diag}(X^T X)$ refers to the matrix where the non-diagonal elements of $X^T X$ are set to 0, and $\tilde{\Gamma} = (\text{diag}(X^T X))^{1/2}$, which exists as $X^T X$ is PSD and therefore has non-negative diagonal elements.

2. Feature Dimensions in CNN

We are going to describe a convolutional neural net using the following pieces:

- CONV3-F denotes a convolutional layer with F different filters, each of size $3 \times 3 \times C$, where C is the depth (i.e. number of channels) of the activations from the previous layer. Padding is 1, and stride is 1.
- POOL2 denotes a 2×2 max-pooling layer with stride 2 (pad 0)
- FLATTEN just turns whatever shape input tensor into a one-dimensional array with the same values in it.
- FC-K denotes a fully-connected layer with K output neurons.

Note: All CONV3-F and FC-K layers have biases as well as weights. **Do not forget the biases when counting parameters.**

We are going to use this network to do inference on a single input. Fill in the missing entries in this table of the size of the activations at each layer, and the number of parameters at each layer. You can/should write your answer as a computation (e.g. $128 \times 128 \times 3$) in the style of the already filled-in entries of the table.

Layer	Number of Parameters	Dimension of Activations
Input	0	$28 \times 28 \times 1$
CONV3-10	Solution: $3 \times 3 \times 1 \times 10 + 10$	$28 \times 28 \times 10$
POOL2	0	$14 \times 14 \times 10$
CONV3-10	$3 \times 3 \times 10 \times 10 + 10$	Solution: $14 \times 14 \times 10$
POOL2	Solution: 0	Solution: $7 \times 7 \times 10$
FLATTEN	0	490
FC-3	Solution: $490 \times 3 + 3$	3